

# 1. Introduction

This document provides the complete system design for a scalable, multi-region URL shortener with real-time analytics.

The system allows users to generate short links, track clicks, referrers, devices, geo-location, and access a dashboard for analytics.

The design covers requirements, architecture diagrams (text), engineering design, data flow, capacity planning, SLOs, and trade-offs.

---

## 2. Requirements Pack

### 2.1 Stakeholder Analysis

#### Primary Stakeholders

- **End Users:** Create and use short URLs, view analytics.
- **Businesses/Marketing Teams:** Require high uptime, fast redirects, accurate analytics.
- **Administrators/Engineers:** Maintain system reliability, monitor performance.

#### Secondary Stakeholders

- **Data Analysts:** Use analytics to study user behavior.
  - **Executives/Product Owners:** Need insights for decision making.
- 

### 2.2 Requirement Prioritization

#### P0 (Must-Have)

- Generate short URLs

- Redirect to original URL
- Track click analytics
- Store URL mappings
- Multi-region high availability

### **P1 (Should-Have)**

- Geo-location tracking
- Device & browser analytics
- Real-time dashboard

### **P2 (Nice-to-Have)**

- Custom short codes
  - QR code generator
  - Authenticated user accounts
- 

## **2.3 Functional Requirements**

1. Users can submit long URLs to generate short codes.
2. System redirects short URL → original URL.
3. The system logs click events:
  - timestamp
  - IP address
  - device & browser
  - referrer
  - country

4. Analytics dashboard displays:
    - total clicks
    - active links
    - CTR
    - top performing links
  5. Multi-region read replicas ensure low-latency redirection.
- 

## 2.4 Non-Functional Requirements

- **Availability:** 99.99% uptime
  - **Latency:** Redirect < 50ms (p95)
  - **Scalability:** Up to millions of redirects/day
  - **Security:** Rate limiting, API keys, HTTPS
  - **Durability:** No lost mappings or analytics
- 

## 2.5 Constraints

- Must operate globally with a multi-region setup.
  - Should minimize database load (read-heavy system).
  - Must support burst traffic during marketing events.
- 

## 2.6 Assumptions

- 80% of traffic are read (redirect).

- Cache will achieve at least 90–95% hit ratio.
  - Users access service from multiple regions globally.
- 

## 3. System Context Diagram (Text Representation)

### Actors & Interactions:

- **User** → accesses website or short link
  - **Web App** → provides UI
  - **API Gateway** → manages secure API routing
  - **URL Shortener Service** → stores and resolves URLs
  - **Analytics Service** → processes click events
  - **Data Stores** → NoSQL DB, Redis, Data Warehouse
  - **CDN/Edge Nodes** → serve cached redirect responses
- 

## 4. Use Case Diagram (Text)

### Use Case 1: Shorten URL

Actor: User

System: Stores mapping and returns short code.

### Use Case 2: Redirect

Actor: Anonymous client user

System: Resolves code → redirects → logs event.

## Use Case 3: View Analytics

Actor: Authenticated user

System: Fetches metrics from analytics DB.

---

## 5. User Stories

- **As a user**, I want to shorten long URLs to share them easily.
  - **As a user**, I want to track how many clicks my short URL gets.
  - **As a business**, I need high uptime and low-latency redirects.
  - **As an admin**, I want to monitor system health.
- 

## 6. Core Sequence Diagrams (Text)

### 6.1 URL Shortening Flow

1. User → submits long URL
  2. API Gateway → forwards to URL Shortener service
  3. Service validates URL
  4. Generates unique short code
  5. Stores in ShortLink DB
  6. Returns short URL
- 

### 6.2 Redirection Flow

1. User clicks short URL
  2. CDN checks cache →
    - if hit → redirect immediately
    - if miss → fetch from DB
  3. Service publishes click event to Queue
  4. System redirects user
  5. Analytics service processes events
- 

### 6.3 Analytics Flow

1. Event Queue receives click event
  2. Analytics Worker extracts device, geo, referrer
  3. Writes aggregated metrics to OLAP DB
  4. Dashboard fetches analytics from analytics DB
- 

## 7. High-Level Architecture Overview

### Core Components

- **API Gateway**
- **URL Shortener Service**
- **Redirection Service**
- **Analytics Ingestion Service**
- **Message Queue (Kafka / SQS / PubSub)**
- **ShortLink DB (NoSQL)**

- **Redis Cache**
- **CDN/Edge Network**
- **Analytics OLAP Database**

## Data Flow

- Shorten URL → Write DB
  - Redirect → Cache → DB fallback
  - Click event → Queue → Analytics DB
- 

# 8. Engineering Notes

## 8.1 Capacity Planning & Sizing

Assume:

- 10 million requests/day
  - 80% redirects = 8M redirect requests
  - Expected cache hit ratio: 95%
  - DB reads: 5% of 8M = ~400k/day
  - Storage per URL:
    - short code + long URL ≈ 100 bytes
    - 10M URLs ≈ 1GB
- 

## 8.2 API Specifications

### POST /shorten

Request: longURL

Response: shortURL

### **GET /{code}**

Redirects to original URL

### **GET /analytics/{code}**

Returns metrics

---

## **8.3 Data Model**

### **ShortLink Table**

- code (primary key)
- longURL
- createdAt
- userId

### **Analytics Event**

- eventId
  - code
  - timestamp
  - IP
  - device
  - country
  - referrer
- 

## **8.4 Consistency Model**

- **Strong consistency** for URL creation
  - **Eventual consistency** for analytics
- 

## 8.5 Caching Strategy

- Redis for code → URL mapping
  - TTL optional
  - LRU eviction
- 

## 8.6 Indexing

- Primary index on code
  - Secondary index: userId (optional)
- 

## 8.7 Rate Limiting

- Throttle:
    - API calls per IP
    - URL creation per minute
  - Prevents abuse and DDoS
- 

## 8.8 Resiliency

- Retries with exponential backoff

- Timeouts on DB and queue calls
  - Circuit breakers for failing services
  - Multi-region failover
- 

## 8.9 Observability

### Logs

- Redirect logs
- Error logs

### Metrics

- QPS
- latency
- cache hit ratio
- queue lag

### Tracing

- API gateway traces
  - DB interactions
- 

## 8.10 Maintenance Plan

- Archive old analytics
- Scheduled DB backups
- Index optimization

- Auto-scaling workers and API servers
- 

## 9. Quality Targets

### 9.1 SLOs

- Redirect latency: **<50ms (p95)**
  - Uptime: **99.99%**
  - Analytics delay: **<5 seconds**
- 

### 9.2 Scalability Plan

- Stateless services → horizontal scaling
  - DB sharding by:
    - shortcode prefix
    - userId (optional)
  - Partitioned analytics storage
- 

### 9.3 Trade-Offs

- **NoSQL vs SQL:** NoSQL chosen for scalability and fast reads.
- **Eventual consistency** for analytics improves throughput.
- **CDN caching** reduces DB load but may require invalidation.

**Shortcode generation:**

- Sequential IDs → fast but predictable
  - Random codes → safer but more collision checks
- 

## 10. Conclusion

This design delivers a globally distributed, highly scalable URL shortener supporting real-time analytics, high availability, and low latency. It is suitable for enterprise-grade workloads and marketing use cases.